

Assignment: Image Upload API with AWS LocalStack, DynamoDB, S3, and Testing

1. Overview

This project implements an image management system using AWS services:

- **S3**: Stores image files.
- **DynamoDB**: Stores image metadata (file name, unique ID, created_at).
- **Lambda Functions**: Handle CRUD operations (upload, list, get, delete).
- **LocalStack**: Emulates AWS services locally for development.
- **Pytest + Moto**: Unit and integration testing.

The API supports:

- Uploading images (POST)
- Listing images (GET)
- Fetching images by ID (GET)
- Deleting images (DELETE)

2. Project Structure

```
monty-cloud-assignment/  
|  
├── Dockerfile.localstack  
├── docker-compose.yml  
├── handler.py          # Lambda handlers for CRUD operations  
└── setup-localstack.sh # Script to setup LocalStack resources
```

```
|— requirements.txt    # Python dependencies (boto3, pytest, moto, etc.)
|— tests/
|   |— test_handler.py # Pytest test cases
|— logs/              # LocalStack logs
|— handler.py         # Mounted handler code inside container
```

3. Lambda Handlers Overview

The Lambda functions perform CRUD operations with the following key considerations:

- **Upload:**
 - Accepts file name, base64-encoded file content, and optional metadata.
 - Generates a unique ID for each upload.
 - Saves the file to S3 and metadata to DynamoDB.
- **List Images:**
 - Scans the DynamoDB table and returns all image metadata.
 - Supports optional filtering by file name or creation date.
- **Get Image:**
 - Retrieves the image metadata from DynamoDB using the unique ID.
 - Fetches the corresponding file content from S3.
- **Delete Image:**
 - Deletes the file from S3 and metadata from DynamoDB.
 - Handles repeated deletions gracefully.
- **Error Handling:**
 - Returns proper HTTP codes: 2xx for success, 4xx for client errors, and 5xx for server errors.

4. Docker Setup

docker-compose.yml:

- Runs LocalStack in a container.
- Exposes the edge port for API requests.
- Mounts local files for handler code, setup script, and tests.
- Includes Docker socket to allow Lambda to run inside LocalStack.
- Installs dependencies (boto3, pytest, moto, awscli-local).

setup-localstack.sh:

- Creates the S3 bucket.
- Creates the DynamoDB table with a unique ID as the primary key.
- Ensures LocalStack resources are ready for Lambda invocation.

Ensure the setup script has execution permissions.

5. Testing with Pytest

- **Fixtures:**
 - Sets up environment variables (`BUCKET_NAME`, `TABLE_NAME`).
 - Mocks AWS resources using Moto (`mock_aws`).
- **Test Coverage:**
 - **Success Cases:** Upload, list, get, delete.

- **Failure Cases:** Invalid HTTP method, missing fields, invalid base64, non-existent IDs.
 - **Edge Cases:** Large payloads, double deletion, DynamoDB or S3 failures.
 - **HTTP Codes:** Validates 2xx, 4xx, and 5xx responses.
 - **Notes on Moto:**
 - Moto's `mock_aws` decorator mocks multiple services in one context.
 - Always specify a region to avoid `NoRegionError`.
-

6. Running the Project

Start LocalStack

- Use Docker Compose to start LocalStack in detached mode.

Invoke Lambda Functions Locally

- Use `docker exec -it localstack sh /app/setup-localstack.sh` to test Lambda functions directly.

Run Tests Inside Container

- Use `docker exec -it localstack pytest -v /app/tests/test_handler.py` to execute test cases inside the LocalStack container.
-

7. Notes & Recommendations

1. Unique IDs:

- Use UUIDs to avoid collisions for parallel requests.

- Guarantees uniqueness for high-concurrency uploads.

2. Error Handling:

- Return clear HTTP 4xx/5xx messages.
- Handle missing files, missing IDs, and invalid payloads.

3. Testing:

- Cover success, failure, edge, and large payload scenarios.
- Always mock AWS services for unit testing to avoid hitting real endpoints.

4. LocalStack Considerations:

- Mount `/var/run/docker.sock` to allow Lambda Docker execution.
- Ensure endpoint URLs are correctly set when calling S3/DynamoDB.

5. Environment Variables:

- Use `BUCKET_NAME`, `TABLE_NAME`, `LOCALSTACK_ENDPOINT` in code and tests.

6. Performance:

- DynamoDB primary key is `id` to ensure efficient lookups.
- Use `scan` carefully; for large datasets consider `query` with indexes.